

Table of Contents

Preface	1
Chapter 1: Getting Started	6
Introduction	6
Installing R with an IDE	7
Getting ready	7
How to do it...	8
Installing a Jupyter Notebook application	9
How to do it...	10
There's more...	11
Starting with the basics of machine learning in R	12
How to do it...	13
How it works...	17
Setting up deep learning tools/packages in R	18
How to do it...	18
Installing MXNet in R	19
Getting ready	19
How to do it...	20
Installing TensorFlow in R	21
Getting ready	21
How to do it...	21
How it works...	22
See also	23
Installing H2O in R	24
Getting ready	24
How to do it...	24
How it works...	26
There's more...	29
Installing all three packages at once using Docker	30
Getting ready	30
How to do it...	31
There's more...	32
Chapter 2: Deep Learning with R	33
Starting with logistic regression	33
Getting ready	34

How to do it...	34
Introducing the dataset	35
Getting ready	35
How to do it...	36
Performing logistic regression using H2O	37
Getting ready	37
How to do it...	38
How it works...	40
See also	41
Performing logistic regression using TensorFlow	42
Getting ready	42
How to do it...	42
How it works...	44
Visualizing TensorFlow graphs	45
Getting ready	45
How to do it...	46
How it works...	49
Starting with multilayer perceptrons	50
Getting ready	50
How to do it...	51
There's more...	52
See also	53
Setting up a neural network using H2O	53
Getting ready	53
How to do it...	54
How it works...	56
Tuning hyper-parameters using grid searches in H2O	57
Getting ready	57
How to do it...	57
How it works...	58
Setting up a neural network using MXNet	59
Getting ready	59
How to do it...	59
How it works...	61
Setting up a neural network using TensorFlow	61
Getting ready	61
How to do it...	62
How it works...	65
There's more...	66

Chapter 3: Convolution Neural Network	68
Introduction	69
Downloading and configuring an image dataset	70
Getting ready	70
How to do it...	70
How it works...	73
See also	74
Learning the architecture of a CNN classifier	74
Getting ready	74
How to do it...	75
How it works...	76
Using functions to initialize weights and biases	77
Getting ready	78
How to do it...	78
How it works...	78
Using functions to create a new convolution layer	78
Getting ready	79
How to do it...	79
How it works...	81
Using functions to create a new convolution layer	82
Getting ready	82
How to do it...	82
How it works...	83
Using functions to flatten the densely connected layer	83
Getting ready	83
How to do it...	84
How it works...	84
Defining placeholder variables	85
Getting ready	85
How to do it...	85
How it works...	86
Creating the first convolution layer	86
Getting ready	86
How to do it...	87
How it works...	87
Creating the second convolution layer	88
Getting ready	88
How to do it...	88
How it works...	89

Flattening the second convolution layer	90
Getting ready	90
How to do it...	91
How it works...	91
Creating the first fully connected layer	91
Getting ready	91
How to do it...	92
How it works...	92
Applying dropout to the first fully connected layer	92
Getting ready	92
How to do it...	92
How it works...	93
Creating the second fully connected layer with dropout	93
Getting ready	93
How to do it...	94
How it works...	94
Applying softmax activation to obtain a predicted class	94
Getting ready	94
How to do it...	95
Defining the cost function used for optimization	95
Getting ready	95
How to do it...	95
How it works...	96
Performing gradient descent cost optimization	96
Getting ready	96
How to do it...	96
Executing the graph in a TensorFlow session	97
Getting ready	97
How to do it...	97
How it works...	98
Evaluating the performance on test data	99
Getting ready	99
How to do it...	99
How it works...	101
Chapter 4: Data Representation Using Autoencoders	104
Introduction	104
Setting up autoencoders	106
Getting ready	107
How to do it...	107

Data normalization	108
Getting ready	108
Visualizing dataset distribution	109
How to do it...	110
How to set up an autoencoder model	111
Running optimization	114
Setting up a regularized autoencoder	115
Getting ready	115
How to do it...	116
How it works...	116
Fine-tuning the parameters of the autoencoder	117
Setting up stacked autoencoders	118
Getting ready	119
How to do it...	120
Setting up denoising autoencoders	121
Getting ready	121
How to do it...	121
Reading the dataset	121
Corrupting data to train	122
Setting up a denoising autoencoder	124
How it works...	126
Building and comparing stochastic encoders and decoders	127
Getting ready	128
How to do it...	129
Setting up a VAE model	130
Output from the VAE autoencoder	134
Learning manifolds from autoencoders	135
How to do it...	135
Setting up principal component analysis	136
Evaluating the sparse decomposition	139
Getting ready	140
How to do it...	140
How it works...	141
Chapter 5: Generative Models in Deep Learning	143
Comparing principal component analysis with the Restricted Boltzmann machine	144
Getting ready	145
How to do it...	145
Setting up a Restricted Boltzmann machine for Bernoulli distribution input	149

Getting ready	150
How to do it...	150
Training a Restricted Boltzmann machine	151
Getting ready	151
Example of a sampling	151
How to do it...	152
Backward or reconstruction phase of RBM	152
Getting ready	153
How to do it...	153
Understanding the contrastive divergence of the reconstruction	154
Getting ready	154
How to do it...	154
How it works...	155
Initializing and starting a new TensorFlow session	155
Getting ready	156
How to do it...	156
How it works...	157
Evaluating the output from an RBM	158
Getting ready	159
How to do it...	159
How it works...	160
Setting up a Restricted Boltzmann machine for Collaborative Filtering	162
Getting ready	162
How to do it...	162
Performing a full run of training an RBM	164
Getting ready	167
How to do it...	167
Setting up a Deep Belief Network	169
Getting ready	171
How to do it...	171
How it works...	174
Implementing a feed-forward backpropagation Neural Network	175
Getting ready	176
How to do it...	176
How it works...	180
Setting up a Deep Restricted Boltzmann Machine	180
Getting ready	181
How to do it...	181
How it works...	186

Chapter 6: Recurrent Neural Networks	188
Setting up a basic Recurrent Neural Network	188
Getting ready	189
How to do it...	189
How it works...	192
Setting up a bidirectional RNN model	193
Getting ready	193
How to do it...	194
Setting up a deep RNN model	197
How to do it...	197
Setting up a Long short-term memory based sequence model	198
How to do it...	198
How it works...	199
Chapter 7: Reinforcement Learning	202
Introduction	202
Setting up a Markov Decision Process	204
Getting ready	204
How to do it...	205
Performing model-based learning	210
How to do it...	212
Performing model-free learning	213
Getting ready	213
How to do it...	215
Chapter 8: Application of Deep Learning in Text Mining	218
Performing preprocessing of textual data and extraction of sentiments	218
How to do it...	219
How it works...	225
Analyzing documents using tf-idf	226
How to do it...	226
How it works...	228
Performing sentiment prediction using LSTM network	229
How to do it...	229
How it works...	233
Application using text2vec examples	233
How to do it...	233
How it works...	236
Chapter 9: Application of Deep Learning to Signal processing	237

Introducing and preprocessing music MIDI files	237
Getting ready	238
How to do it...	239
Building an RBM model	239
Getting ready	239
How to do it...	240
Generating new music notes	242
How to do it...	242
Chapter 10: Transfer Learning	243
Introduction	243
Illustrating the use of a pretrained model	245
Getting ready	246
How to do it...	246
Setting up the Transfer Learning model	249
Getting ready	249
How to do it...	249
Building an image classification model	252
Getting ready	252
How to do it...	252
Training a deep learning model on a GPU	256
Getting ready	256
How to do it...	256
Comparing performance using CPU and GPU	257
Getting ready	258
How to do it...	258
There's more...	260
See also	260
Index	261