

Digital Filtering

Prof Gelman

Introduction

- Digital filtering is an important function that can be implemented in the DSP unit
- We use the term *filter* to describe a **linear** system used to perform frequency-selective filtering
- The mathematical foundation of filtering is **the convolution**

Analogue and Digital Filters








- ❑ Digital filters are now extremely cheap and their use is usually preferred over analogue filters for the following main reasons:
 - Digital filters **are programmable**
 - They do not suffer from the **aging distortion**
 - It is extremely difficult to realize RLC analogue filters with low passband edge frequencies and good stopband attenuation; however, it is possible to realize digital filters with very low passband edge frequencies which provide good stopband attenuation

Filter Classification

Filters are usually classified according to their frequency-domain characteristics as


- **lowpass**, passes only low frequencies
- **highpass**, passes only high frequencies
- **bandpass**, passes only frequencies within the selected band
- **bandstop** (or band-elimination filters), eliminates frequencies within the selected band
- The ideal filters have a **constant-gain** (usually taken as **unity-gain**) in their passband and **zero gain** in their stopband

Filter Classification



The Ideal Filter

The ideal lowpass filter is shown below



The Ideal Filter

Ideal filters are physically unrealizable because


- the frequency response of the filter *cannot be zero* except at a finite set of points in the frequency range
- frequency response cannot have an *infinitely sharp cutoff* from passband to stopband
- In addition, the frequency response characteristics possessed by ideal filters *are not absolutely necessary* in most practical applications

The Ideal and Real Filters

- ❑ In particular, it is not necessary that the magnitude of the frequency response **to be constant in the entire passband**
- ❑ Similarly, it is not necessary for the magnitude **to be zero in the entire stopband**
- ❑ A small amount of ripples is usually tolerable

Basic Lowpass Filter

The normalized magnitude frequency response of the basic **non-ideal** lowpass filter is as follows



Lowpass Filter: the Passband Frequency

- The passband edge frequency $\overline{\omega}_p$ defines the **edge of the passband**
- If there is ripple in the passband, its value is denoted as δ_1
- Thus, we require that in the passband the magnitude approximates **unity** with an error of $\pm \delta_1$, e.g.

$$1 - \delta_1 \leq |H(\omega)| \leq 1 + \delta_1 \text{ for } |\omega| \leq \overline{\omega}_p$$

Lowpass Filter: the Stopband Frequency

- The stopband edge frequency $\overline{\omega}_s$ denotes **the beginning of the stopband**
- The ripple in the stopband is denoted as δ_2
- Thus, we require that in the stopband the magnitude approximates **zero** with an error of $\pm \delta_2$, e.g.

$$|H(\omega)| \leq \delta_2 \text{ for } |\omega| \geq \overline{\omega}_s$$

Lowpass Filter: the Transition Band

- The transition of the frequency response from passband to stopband defines the *transition band* of the filter
- Thus, the width of the transition band is $\overline{\omega}_s - \overline{\omega}_p$
- Since all filter design techniques are developed in terms of *normalized frequencies*, the specified frequencies need to be normalized by sampling frequency f_s

$$\omega_p = \frac{\overline{\omega}_p}{f_s}$$

$$\omega_s = \frac{\overline{\omega}_s}{f_s}$$

Lowpass Filter: a Specification

- *In any filter design, we need to specify $\delta_1, \delta_2, \omega_p$ and ω_s*
- **Based on this specification, we can design a filter**

Lowpass Filter


Filter Specification

δ_1 Peak Passband Ripple: $-20\log_{10}(1-\delta_1)$

δ_2 Peak Stopband Ripple $-20\log_{10}(\delta_2)$

ω_p Passband edge frequency

ω_s Stopband edge frequency



The Transfer Function

- The *transfer function* of a discrete-time filter is defined as

$$H(z) = \frac{Y(z)}{X(z)}$$

where $Y(z)$ denotes the z -transform of the filter output signal, and $X(z)$ denotes the z -transform of the filter input signal

- The transfer function is the z -transform of the **filter impulse response function** $h(n)$
- **Filter impulse response function** is a filter response to the unit impulse excitation

The Frequency Response Function


- We can obtain the **frequency response function** (generally, we shorten this to the *frequency response*) of the filter by evaluating the transfer function on the unit circle

- Thus

$$H(\omega) = H(z) \Big|_{z=e^{i\omega}}$$

- Frequency response is the **ratio of the complex Fourier transform of filter output to the Fourier transform of filter input**

Lowpass Filter



Linear Phase Filters

- **Let's consider a filter with a linear phase in the frequency response:**

$$H(\omega) = \begin{cases} Ae^{-i\omega\alpha} & \omega_1 \leq \omega \leq \omega_2 \\ 0, & \text{otherwise} \end{cases}$$

where A and α are constants


- **The Fourier transform of the output of the filter is**

$$Y(\omega) = X(\omega)H(\omega) = AX(\omega)e^{-i\omega\alpha}$$

- **Applying the scaling and time-shifting properties of the Fourier transform, we obtain the time-domain output**

$$y(n) = Ax(n - \alpha)$$

Linear Phase Filter



Linear Phase Filters

- The filter output is simply a delayed and amplitude-scaled version of the input signal
- A pure delay is usually tolerable and is not considered a distortion of the signal. Neither is amplitude scaling
- Therefore, ideal filters **have a constant magnitude characteristic and a linear phase characteristic within their passband**
- In all cases, such filters are not physically realizable, but serve as a mathematical idealisation of practical filters

Filter Representation

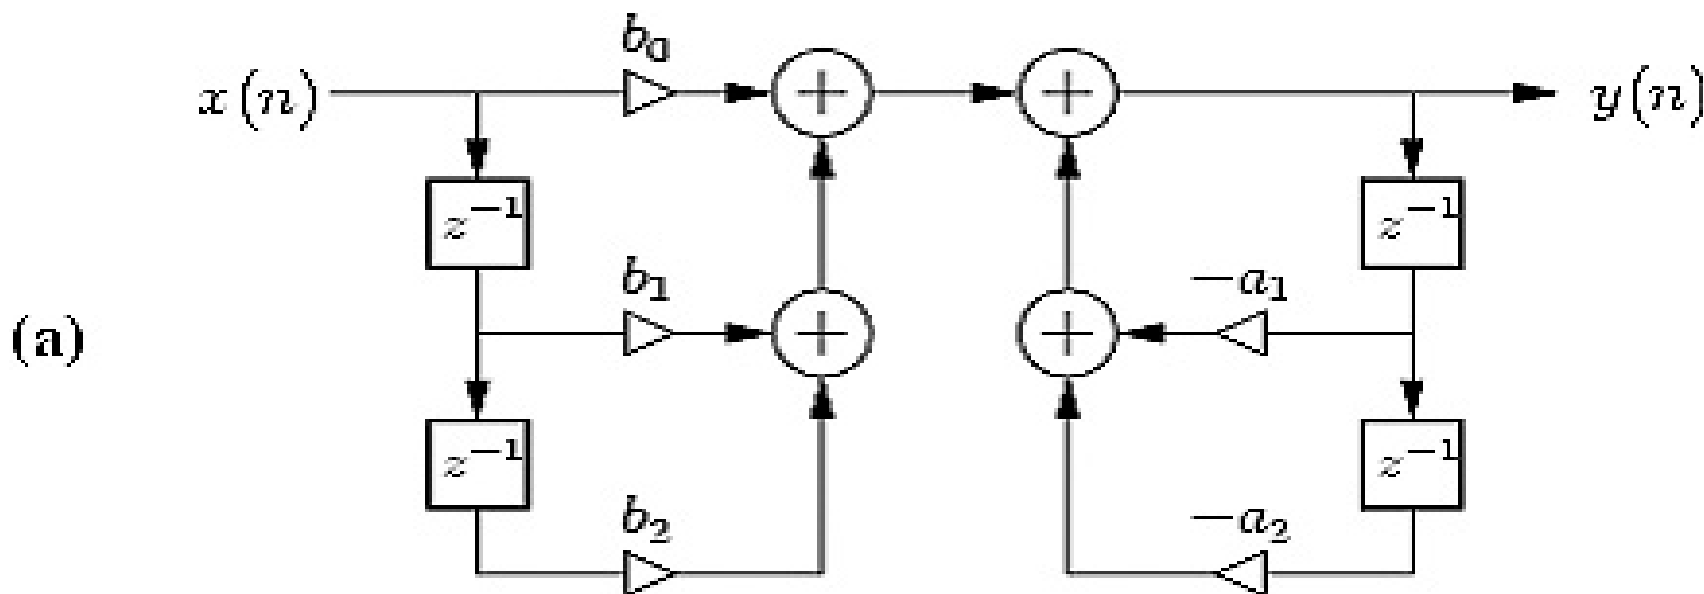
- A digital filter can be described in two main ways
- First, a digital filter can be described as a mathematical algorithm or **constant-coefficient difference** equation, e. g.

$$y(n) = b_0x(n) + b_1x(n-1) + b_2x(n-2) - a_1y(n-1) - a_2y(n-2)$$

- The properties of this particular filter will be entirely described by the constants a_1, a_2, b_0, b_1, b_2

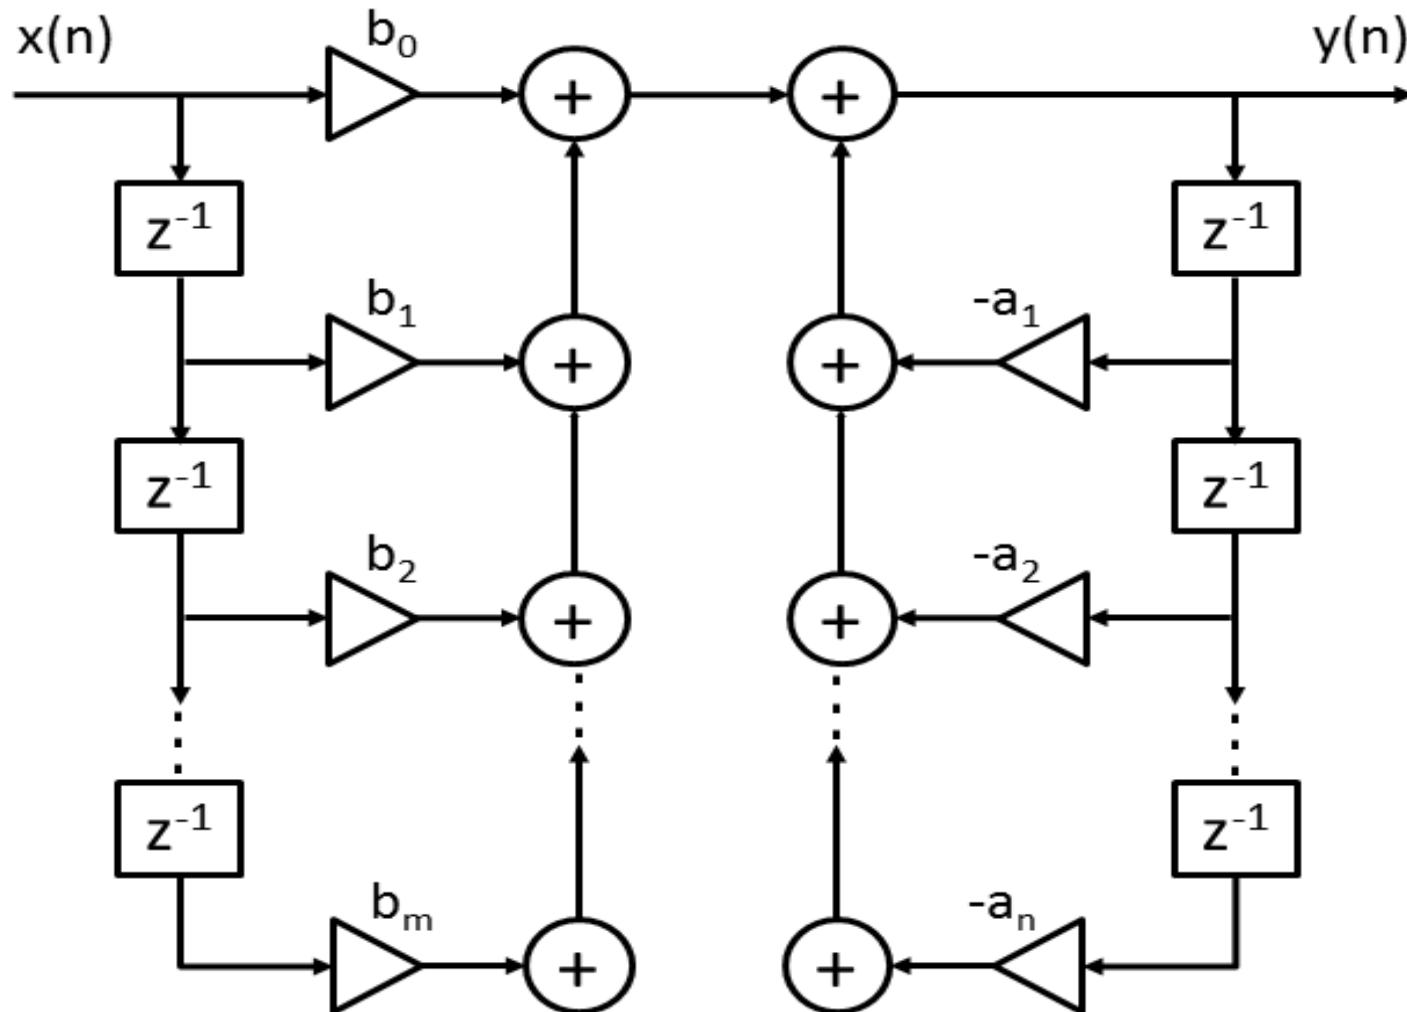
Filter Representation

- The filter can be described graphically as a block diagram




- This diagram represents the above difference equation (a box labeled z^{-1} denotes a one-sample delay in time)

Filter Representation: Direct Structure



Filter Representation



- This structure is known as a *direct structure* of a digital filter and it has two main sections: *feedforward and feedback*
- The input signal is applied directly to the **feedforward** section
- The overall output is formed by adding together the outputs from **feedforward** and **feedback** sections
- “Older” outputs are held in a shift register of the feedback section

The Transfer Function

$$y(n) = b_0x(n) + b_1x(n-1) + b_2x(n-2) - a_1y(n-1) - a_2y(n-2)$$

- **Taking the z-transform of both sides of difference equation and using the delay theorem, it is resulted in:**

$$Y(z) = Z[y(n)] = b_0X(z) + b_1z^{-1}X(z) + b_2z^{-2}X(z) - a_1z^{-1}Y(z) - a_2z^{-2}Y(z)$$

- **Collecting terms, it is resulted in:**

$$H(z) = \frac{b_0 + b_1z^{-1} + b_2z^{-2}}{1 + a_1z^{-1} + a_2z^{-2}}$$

Recursive and Non-Recursive Realizations

- The general form of the filter difference equation is given by

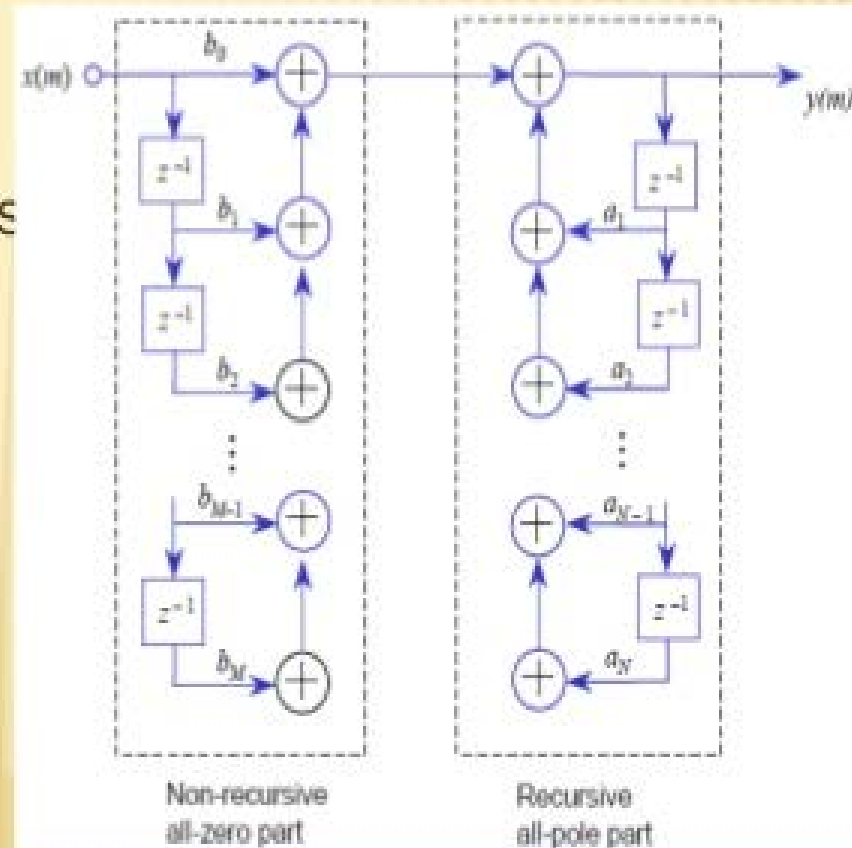
$$y(n) = \sum_{k=0}^{M-1} b_k x(n-k) - \sum_{k=1}^N a_k y(n-k)$$

- This is a *recursive* realization, e.g. the current output is a function of past and present **inputs and outputs**
- For a *non-recursive* realization, the current output is a function only of past and present **inputs**:

$$y(n) = \sum_{k=0}^{M-1} b_k x(n-k)$$

RECURSIVE AND NON RECURSIVE FILTERS

- A recursive filter has feedback from output to input, and in general its output is a function of the previous output samples and the present and past input samples



The Transfer Functions for Recursive and Non-Recursive Realizations

- Taking z-transform of both sides of recursive and non-recursive equations leads to the general forms of the transfer function for **recursive** and **non-recursive** realization respectively


$$H(z) = \frac{\sum_{k=0}^{M-1} b_k z^{-k}}{1 + \sum_{k=1}^N a_k z^{-k}}$$

$$H(z) = \sum_{k=0}^{M-1} b_k z^{-k}$$


Poles and Zeros of the Transfer Function

- The location of poles and zeros affects the frequency response characteristics of filters
- The basic principle underlying the pole-zero placement is to locate poles near points of the unit circle corresponding to frequencies to be emphasized, and to place zeros near the frequencies to be deemphasized
- The position of the poles is an indicator of **filter stability**
- A filter will be stable if all its poles lie inside the *unit circle* in the z-plane; pole modulus should be less than unity


Poles of Stable Filter



Stable and Unstable Regions in Z-Plane




Finite Impulse Response (FIR) Filter

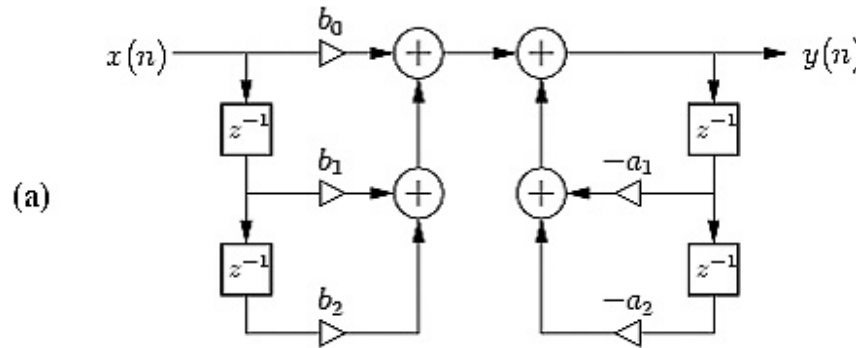


- There is **no feedback section**, all feedback coefficients are zero
- Apply a unit impulse: all samples after it are zero.
- At time zero the output will be b_0 . We then clock the filter which causes the single sample to shift one memory unit down. The output of the filter is now b_1 . We keep clocking the filter until the single impulse reaches the end of the feedforward section, in which case the output will be b_2 . If we clock the filter once more, the output will be zero.
- So, the output is the impulse response (b_0, b_1, b_2)
- The impulse response is **finite** and hence the filter is the **finite impulse response (FIR)**

Finite Impulse Response

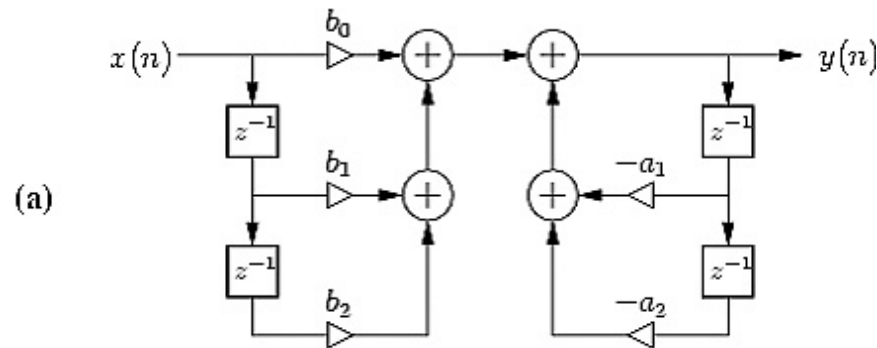


Infinite Impulse Response (IIR) Filter



- For simplicity, let's consider that there is a *feedback* section and *no feedforward section*, i.e. all the coefficients b_i are zero except b_0 which is 1.
- Apply a unit impulse: the first output at $n = 1$ will be 1
We then clock the filter. i. e. $n = 2$
- Although the new input is zero, the previous output value of 1 is sitting at the output of the first delay in the feedback section
Thus, the second output will be $-a_1$

Infinite Impulse Response (IIR) Filter



- We then clock the filter once more, $n = 3$

The previous outputs circulate round the feedback section, i.e.

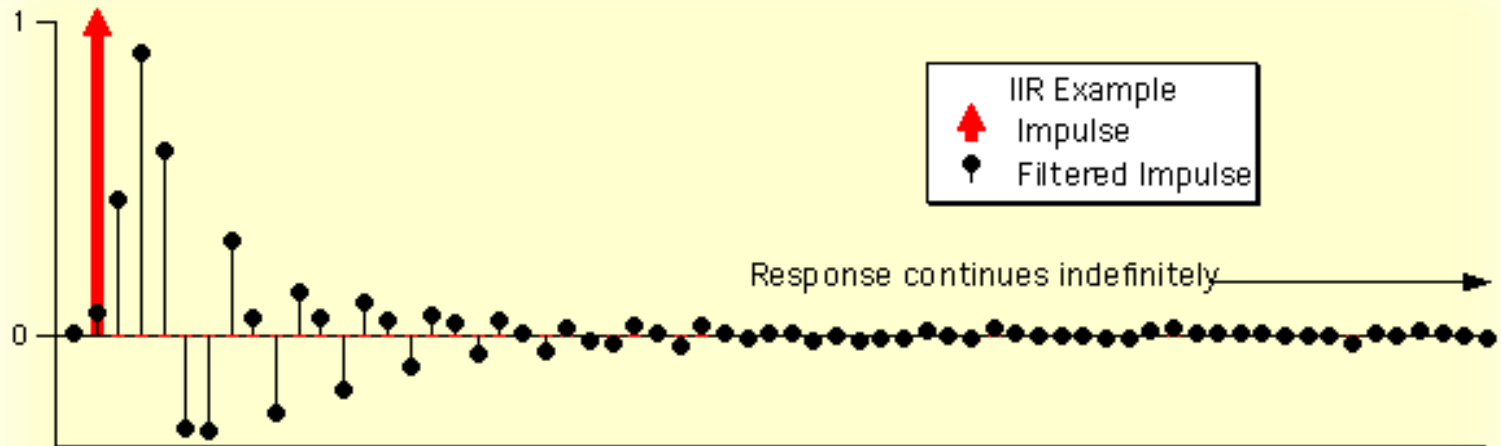
$$y(3) = -a_1 y(2) - a_2 y(1) = a_1^2 - a_2$$

- The next output will be

$$y(4) = 2a_1 a_2 - a_1^3$$

- If we continue to clock the filter, we will continue to get an output as the data circulate round and round the feedback section. **If the filter is stable, the output will decay towards zero but never quite**

Infinite Impulse Response



Infinite Impulse Response (IIR) Filter

The impulse response **is *infinite*** and the filter is the *infinite impulse response* (IIR)

FIR Filters

FIR Filters: Advantages

Among the main *advantages* of FIR filters are:

- FIR filters with exactly *linear phase* can be easily designed
Almost all MATLAB design functions for FIR filters design *linear phase filters only*

Therefore, FIR linear phase filters are important for speech processing and data transmission

- **FIR filters are always stable**
- The filter start up transients have a low duration.

FIR Filters: Equations and the Convolution

- FIR filter is described by the following difference equation:

$$y(n) = \sum_{k=0}^{M-1} b_k x(n-k)$$

- or equivalently by the **convolution**

$$y(n) = \sum_{k=0}^{M-1} h(k)x(n-k) = \sum_{k=0}^{M-1} x(k)h(n-k)$$

- This structure requires $M-1$ memory locations for storing the previous inputs, and has a complexity of M multiplications and $M-1$ additions per output point
- FIR filter defines a weighted running average of M samples

FIR Filters: the Convolution

The parameter M is *order* of the FIR filter

- The process of computing the convolution involves the following four steps:
 - folding
 - shifting
 - multiplication
 - summation

The Transfer Function and Stability

- The transfer function of FIR filter is

$$H(z) = \sum_{k=0}^{M-1} b_k z^{-k}$$

- Now the poles and zeros of this polynomial are identified by expressing the above equation in powers of z which are all **greater than or equal to zero**:

$$H(z) = \frac{b_0 z^{M-1} + b_1 z^{M-2} + \dots + b_{M-1} z^0}{z^{M-1}}$$

- All the denominator poles are at the origin in the z -plane and hence FIR filters are ***unconditionally stable*** as the poles cannot be placed outside the unit circle
- Thus, the frequency response is controlled by the positions of the *numerator zeros*

The Frequency Response

- The frequency response of the FIR filter is obtained directly by replacing z with $\exp(i\omega)$:

$$H(\omega) = \sum_{k=0}^{M-1} b_k \exp(-i\omega k) = \sum_{k=0}^{M-1} h(k) \exp(-i\omega k)$$

- Thus, given a set of **filter weights**, we can evaluate the **complex** frequency response directly
- The above equation is a Fourier series (i.e. DFT) of the frequency response, which is periodic function of ω with period 2π
- Hence, the filter coefficients may be calculated by integrating in the frequency domain over the period of the frequency response (**IDFT**)

$$b_k(n) = \frac{1}{2\pi} \int_{-\pi}^{\pi} H(\omega) \exp(i\omega n) d\omega$$

Computer-Aided Design of Linear-Phase FIR Filters

- Optimized design of linear-phase filters can be achieved by iterative application of **DFT/IDFT** processing
- At the design stage, one has a **desired ideal frequency response: i.e. without ripples in passband and stopband and maximum filter order**
- An optimum design criterion is used in the sense that the **weighted approximation error** between the *desired frequency response* and the *actual frequency response* is spread evenly across the passband and evenly across the stopband of the filter
- In this method, one starts with a set of **filter coefficients** and perform a DFT to get the **actual frequency response** $H(\omega)$

Computer-Aided Design of Linear-Phase FIR Filters

- The filter performance limits are then imposed *on parts of the actual response* $H(\omega)$ *which deviate from desired frequency response* and then the *updated actual response* $H(\omega)$ is given for IDFT to yield an *updated* set of filter coefficients
- This process will result in an updated filter coefficients normally having *component number beyond the maximum filter order permitted for this design*
- *These additional values are simply truncated and then the iterative process is repeated* to get a new actual frequency response $H(\omega)$, which is again compared with the desired frequency response
- This optimization technique is guaranteed to converge

Computer-Aided Design of Linear-Phase FIR Filters

- The objective of this technique is to determine *iteratively* the filter coefficients so that the difference (i.e. weighted error function) between the *desired frequency response*, and *actual frequency response* for all frequencies is minimized
- *The Parks-McClellan algorithm* which reduces significantly the filter complexity compared with window design, uses a minimum weighted Chebyshev error to approximate the desired frequency response $H_d(\omega)$ by iteration:

$$\min \left\{ \max |L(\omega)[H_d(\omega) - H_a(\omega)]| \right\}$$

Computer-Aided Design of Linear-Phase FIR Filters

- The set of filter coefficients, b_k , which *minimizes* the *maximum* error between the desired frequency response, and actual frequency response provides directly **the optimal filter design**
- The positive weighting function $L(\omega)$, allows the designer to *emphasize* some areas of the frequency response more than others
- The minimization is performed *iteratively* using a computer program
- The resulting filter have **ripples** in both the passband and the stopband

IIR Filters

Infinite Impulse Response (IIR) Filters

- *FIR filters is not the most general class of filters*
- The most general class that can be implemented with a finite amount of computation is obtained when the output is formed *not only from the input*, but also from previously computed outputs
- *The primary advantage of IIR filters over FIR filters is that they typically meet a given set of specifications with a much lower filter order than a corresponding FIR filter*

IIR Filters: Equation

- An IIR filter is described by the difference equation:

$$y(n) = \sum_{k=0}^{M-1} b_k x(n-k) - \sum_{k=1}^N a_k y(n-k)$$

- The coefficients a_k are called **feedback coefficients**, and the coefficients b_k are called **the feedforward coefficients**.
- The number N of **feedback terms** is *the order of an IIR filter*
- *In most cases, especially digital filters derived from analog filters, $M \leq N$ in order to prevent infinite gain at infinite frequency*

IIR Filters: the Transfer Function

The general form of the transfer function of IIR filters is

$$H(z) = \frac{\sum_{k=0}^{M-1} b_k z^{-k}}{1 + \sum_{k=1}^N a_k z^{-k}}$$

Linearity and Time Invariance

Although the feedback terms make the proof of linearity and time invariance more complicated than the FIR case, we can easily show that:

- the **principle of superposition** (e.g. *filter linearity*) will hold because the difference equation involves only linear combinations of the input and output samples
- ***time invariance*** of the IIR filters also holds

IIR Filter Design: Analogue Prototyping

- The *most popular* technique for designing IIR filters is based on **converting an analogue filter into digital filter**
- An analogue filter can be described by its transfer function (with Laplace transform):

$$H_a(s) = \frac{\sum_{k=0}^M \beta_k s^k}{\sum_{k=0}^N \alpha_k s^k}$$

where α_k and β_k are the filter coefficients

- *No single transform exists to perfectly map $H_a(s)$ to $H(z)$*

IIR Filter Design: Analogue Prototyping

- In the design of IIR filters by analogue prototyping, we shall specify the desired filter characteristics for the *magnitude frequency response only* and *accept the phase response* that is obtained from the design methodology
- The classic IIR filter design by analogue prototyping can be realized by the following two approaches

IIR Filter Design: Analogue Prototyping

Approach 1

- Find an **basic analogue lowpass filter** and transform this "prototype" filter to the desired frequency band configuration in the *analogue domain*
- Transform the filter to the digital domain

Approach 2

- Find an **basic analogue lowpass filter** and transform this "prototype" filter to the *digital domain*
- Transform digital lowpass filter to the desired frequency band configuration in the **digital domain**

IIR Filter Design: Analogue Prototyping

- **Thus, approach 1 is to perform the frequency transformation in the analogue domain and then to convert the analogue filter into a corresponding digital filter**
- **The approach 2 is first to convert the analogue lowpass filter into a lowpass digital filter and then to transform the lowpass digital filter into the desired digital filter by a digital frequency transformation**
- *In general, these two approaches yield different results*

The Butterworth Analogue Filters


- The Butterworth filters are characterized by the property that the magnitude response is *maximally flat in the passband*
- Another property is that the approximation is *monotonic* in the passband and the stopband
- The squared magnitude transfer function is

$$|H(\Omega)|^2 = \frac{1}{1 + \left(\frac{\Omega}{\Omega_p}\right)^{2N}}$$

where N is the order of the filter, Ω_p is the passband edge frequency

The Butterworth Analogue Filters

- As the order increases, the magnitude become sharper
- They remain closer to unity over more of passband and become close to zero rapidly on the stopband, although the magnitude at the passband edge frequency is always **-3dB**



The Bessel Analogue Filters

- An important characteristic of Bessel analogue filters is the ***linear-phase response over the passband of the filter***; thus, filtered signals maintain their wave shapes in the passband frequency range
- However, we should emphasize that the linear-phase characteristics of the analogue filter ***are destroyed in the process of converting the filter into the digital domain***
- Therefore, digital Bessel filters *do not have* this property

The Bessel Analogue Filters

- Difficulty with Bessel filters is that unlike Butterworth filters the **passband edge frequency varies with the filter order**
- Bessel filters generally require a higher filter order than other filters for **satisfactory stopband attenuation**

The Chebyshev Analogue Filters


- Chebyshev filters are characterized by the property that over a *prescribed frequency band* the approximation error is minimized
- The magnitude error is, in fact, equiripple over the frequency band
- Depending on *whether the band of frequencies* over which the error is minimized is the passband or the stopband, the filter designs are called **Type I or Type II**

The Chebyshev Analogue Filter Type I

- Chebyshev Type I filters minimize the absolute difference between the ideal and actual frequency response over the *entire passband* by incorporating an *equal ripple in the passband*
- Thus, these filters exhibit *equiripple passband behaviour* and monotonic stopband response

The Chebyshev Analogue Filter Type I

- The typical magnitude response for Chebyshev Type I filter



- The optimality property that Type I Chebyshev filters satisfy is there is **no better filter** with equal or better performance in both the passband and stopband

The Chebyshev Analogue Filter Type II


- Type II filters minimize the absolute difference between the ideal and actual frequency response over the **entire stopband** by incorporating an **equal ripple in the stopband**
- These filters exhibit **monotonic** behaviour in the passband and *equiripple behaviour in the stopband*
- The squared-magnitude response of a filter can be expressed as

$$|H(\Omega)|^2 = \frac{1}{1 + \delta_1^2 [T_N(\Omega_r / \Omega_p) / T_N(\Omega_r / \Omega)]^2}$$

where Ω_r is the lowest frequency at which the stopband loss attains a prescribed value

The Chebyshev Analogue Filter Type II

- Typical magnitude response for Chebyshev Type II filter




Elliptic Filters

- Elliptic filters are characterized by a magnitude response that *is equiripple in both the passband and the stopband*
- It can be shown that elliptic filters **are optimum** in the sense that for a given order and for given ripple specifications no other filter achieves a faster transition between the passband and stopband, e.g. they have *the smallest transition band*
- They generally meet filter requirements with the *lowest order* of any supported filter type

Elliptic Filters

The typical magnitude response for elliptic filters



Elliptic Filters: the Phase Response

- The phase response of elliptic filters **is more nonlinear in the passband** than a comparable Butterworth or a Chebyshev filter, especially near the band edge
- In view of the optimality of elliptic filters, one important reason that other types of filters might be preferable in some applications is that they possess **better phase response characteristics**

Comparison of IIR and FIR Filters

- The choice between a FIR filter and an IIR filter depends upon the relative weight that one attaches to the advantages and disadvantages of each type of filter
- If we put aside phase consideration (IIR filters have a nonlinear phase) it is generally true that a given magnitude frequency response specification will be met **more efficiently** with an IIR filter
- It has been shown that for most practical filter specifications, the ratio N_{FIR} / N_{IIR} is typically of the order **of tens**, where N_{FIR} and N_{IIR} are the numbers of multiplications per output sample for an FIR and IIR filters respectively

Comparison of IIR and FIR Filters

- However, IIR design disregards the *phase response* of the filter
- In contrast, FIR filters can have **precisely linear phase**
- In addition, *they always stable*
- In many applications, the linearity of phase response is not an issue
- However, in many cases, the linear phase available with an FIR filter may be **well worth the extra cost**

Comparison of IIR and FIR Filters

- IIR filters have the advantage that they can be designed using **close-form design expressions**
- Most of the other FIR design methods **are iterative procedures**, requiring rather powerful computational facilities for their implementation

Comparison of IIR and FIR Filters

- It has been found that the most favorable conditions for an FIR design are large values of passband ripple, small values of stopband ripple and large transition bands
- In contrast, it is often possible to design frequency selective IIR filters using simple calculations and tables of analogue filter design parameters
- If we consider specific filters, *elliptic IIR filters* are generally more efficient in achieving given specification on the *magnitude response* than optimum FIR filters
- However, its phase response will be very nonlinear (especially at the band edge)